



Word embeddings

Radoslav Neychev, neychev@bigdatateam.org

ML Instructor, [BigDataTeam](#)

Assistant Professor, MIPT

[LinkedIn profile](#)

09.12.2020, online



MEGAFON

- ▶ NLP: introduction
- ▶ Text Preprocessing
- ▶ Feature Extraction: classical approach
 - ▶ Bag-of-Words
 - ▶ Bag-of-Ngramms
 - ▶ TF-IDF
- ▶ Word Embeddings



Natural Language Processing: Introduction



Sentiment analysis



Spam filtering




Fake news detection



Topic prediction




#hashtag prediction



Text classification tasks



Example: sentiment analysis





Discrete labels:

- ▷ Binary
 - spam filtering, sentiment analysis
- ▷ Multi-class
 - categorization of items by its description
- ▷ Multi-label
 - #hashtag prediction

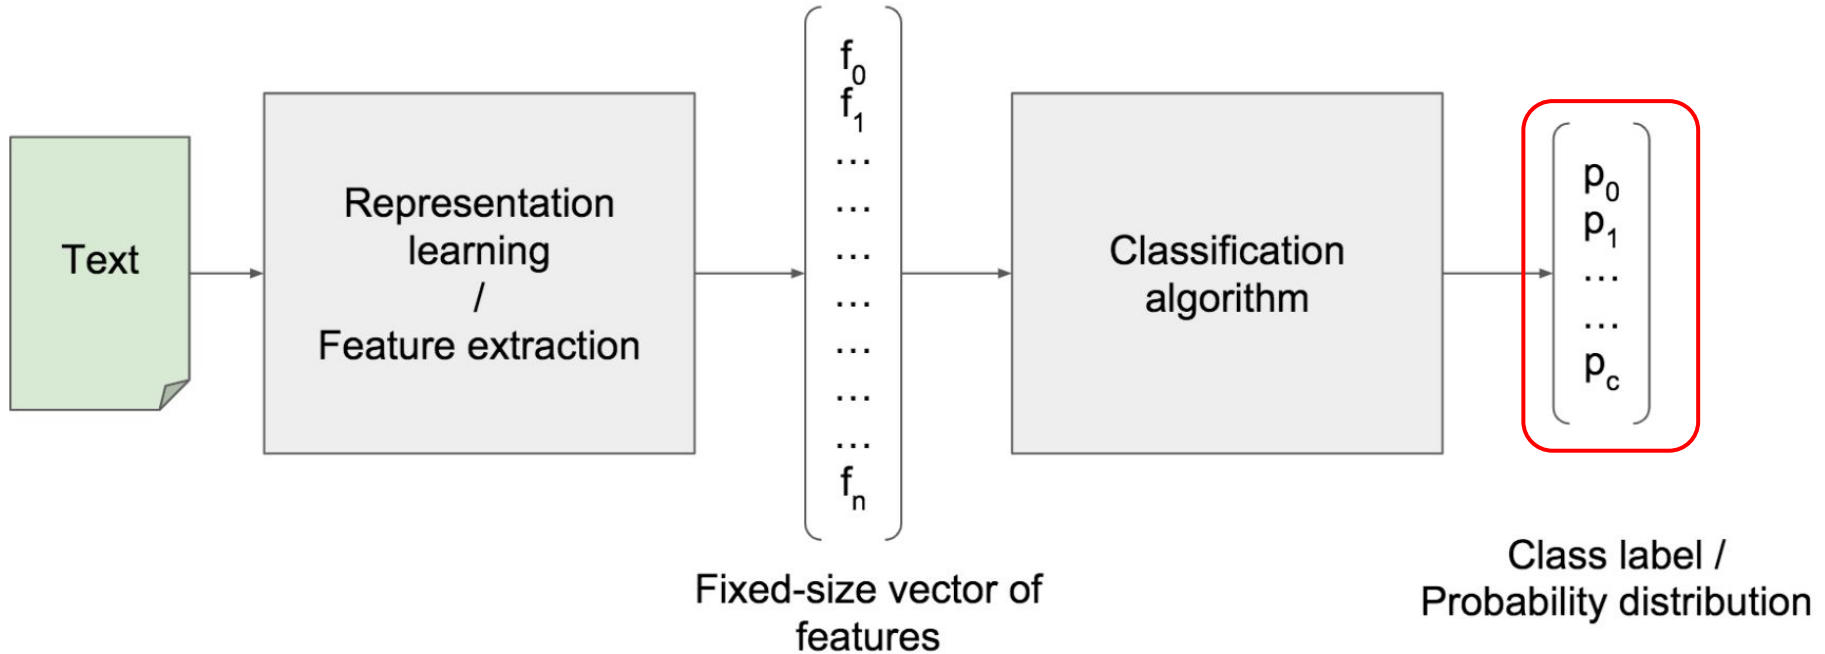


Continuous labels:

- ▷ Predict product price by its description




Text classification in general



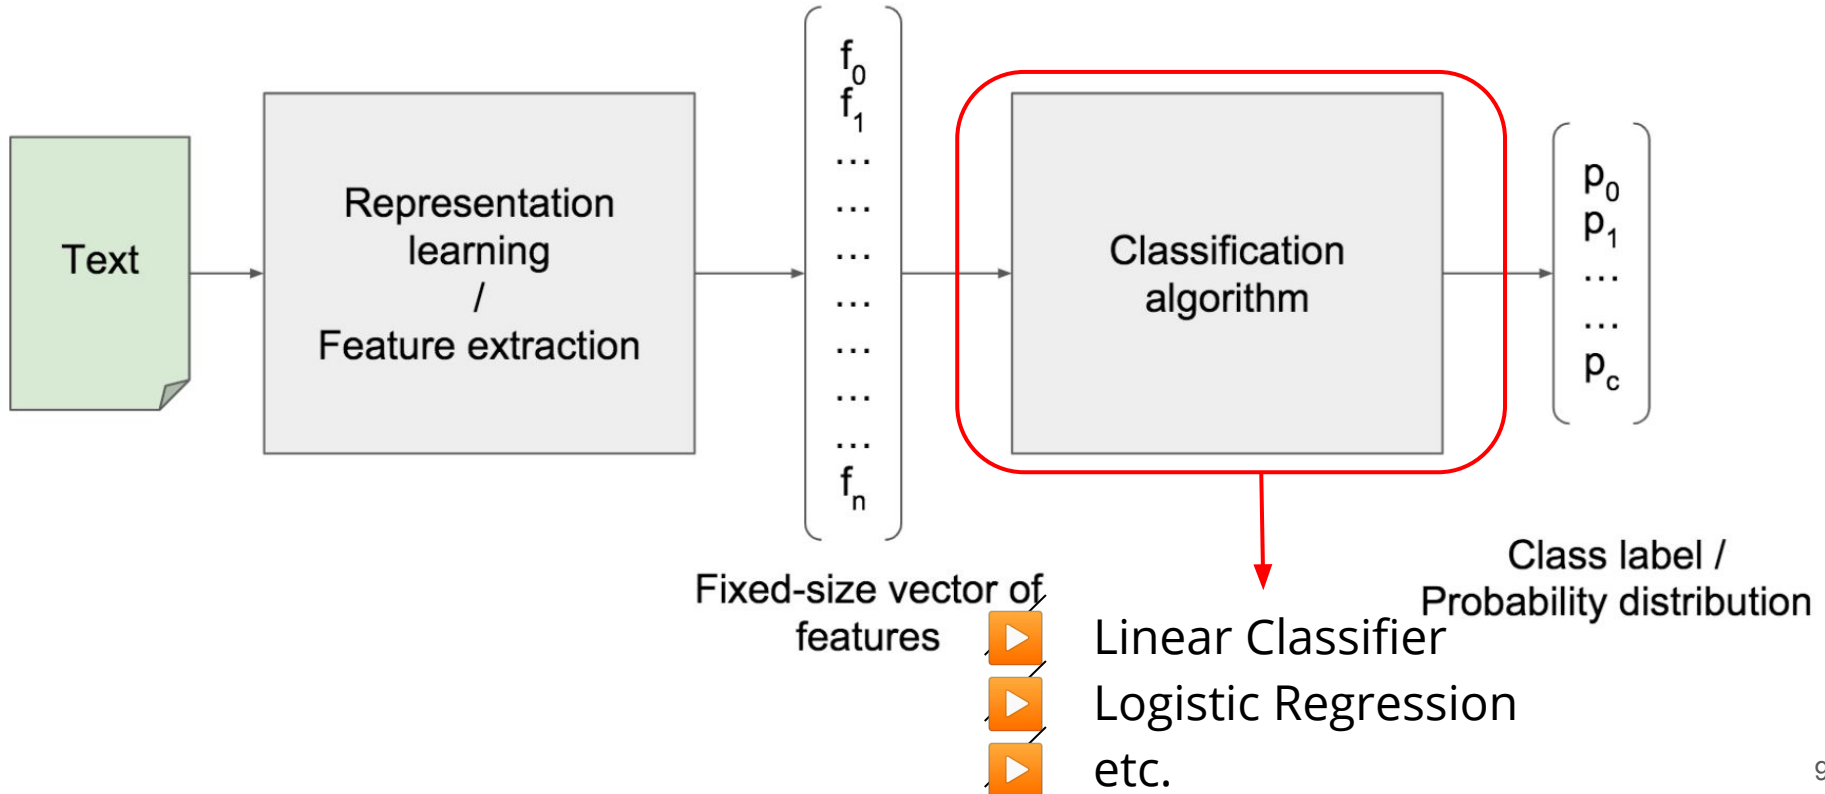


Text classification in general






Text classification in general





Text classification in general






Feature extraction



- ▶ Tokenization: split the input into tokens





the dog is on the table

0	0	1	1	0	1	1	1
are	cat	dog	is	now	on	table	the



Problems:

- ▶ No information about words order
- ▶ Word vectors are huge and very sparse
- ▶ Word vectors are not normalized
- ▶ Same words can take different forms



Text Preprocessing




Token normalization

Dog, dogs → dog

Bark, barks → bark


▶ Token normalization:

- ▶ **Stemming:** removing and replacing suffixes to get to the root of the word (**stem**)



▶ Token normalization:

- ▶ **Stemming:** removing and replacing suffixes to get to the root of the word (**stem**)
- ▶ **Lemmatization:** to get base or dictionary form of a word (**lemma**)





Porter stemmer

- ▶ Published in 1979
- ▶ Base starting option

Snowball stemmer (Porter 2)

- ▶ Based on Porter
- ▶ More aggressive
- ▶ Most popular option now

Lancaster stemmer

- ▶ Published in 1990
- ▶ The most aggressive
- ▶ Easy adding of your own rules



▶ Porter's stemmer:

▷ **Heuristics, applied one-by-one:**


- SSES - SS (dresses - dress)
- IES - I (ponies - poni)
- S - <empty> (dogs - dog)

▷ **What's wrong?**


- **Overstemming and understemming**



Overstemming



Understemming



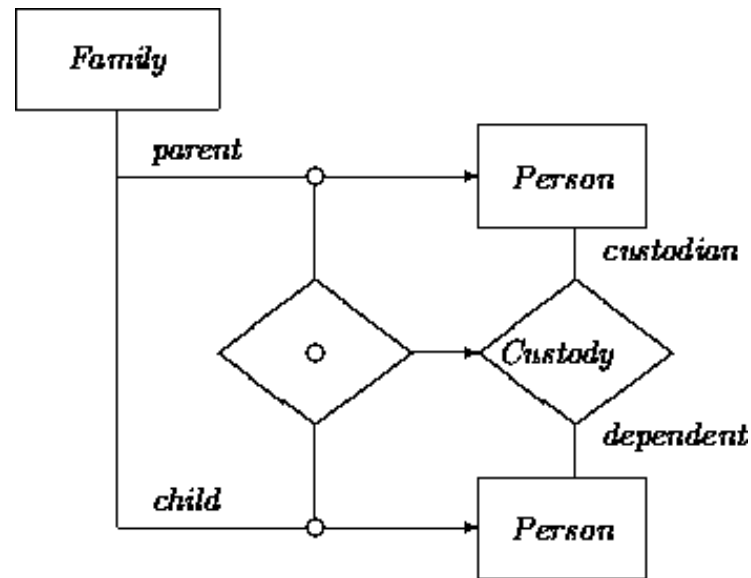
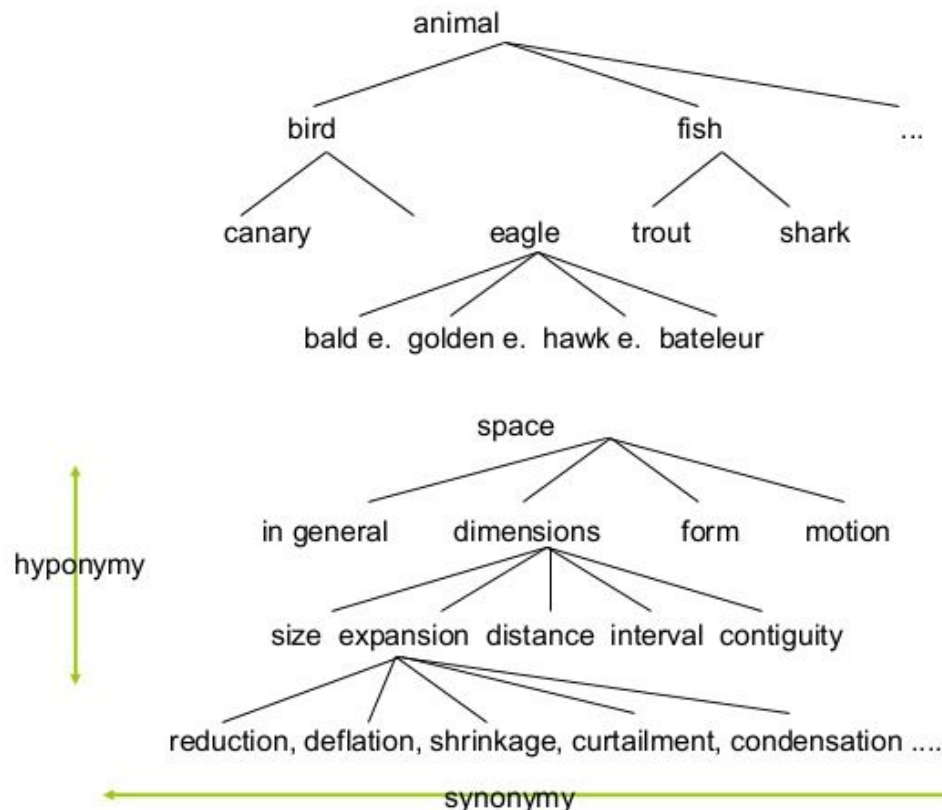


Lemmatizer from NLTK:

- ▷ Tries to resolve word to its dictionary form
- ▷ Based on **WordNet** database
- ▷ For the best results feed part-of-speech tagger



BTW, what is WordNet?



3/27



Handful tools for preprocessing



NLTK

- ▷ `nltk.stem.SnowballStemmer`
- ▷ `nltk.stem.PorterStemmer`
- ▷ `nltk.stem.WordNetLemmatizer`
- ▷ `nltk.corpus.stopwords`



BeautifulSoup (for parsing HTML)



Regular Expressions (import re)



Pymorphy2



- ▶ Capital Letters
- ▶ Punctuation
- ▶ Contractions (e.g, etc.)
- ▶ Numbers (dates, ids, page numbers)
- ▶ Stop-words ("the", "is", etc.)
- ▶ Tags



How to improve BOW?

- ▶ Use n-gramms instead of words!

The brown dog plays
with a little cat



The brown

brown dog

dog plays

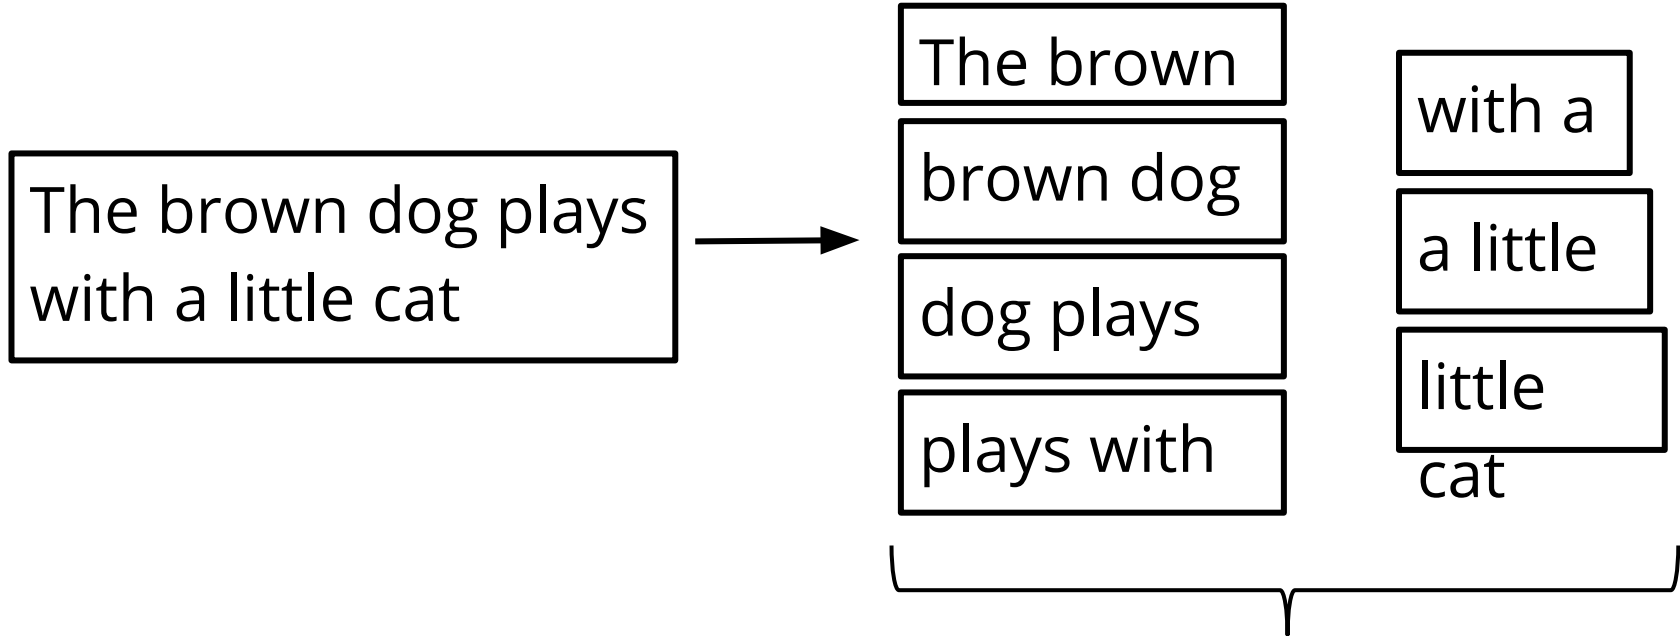
plays with

with a

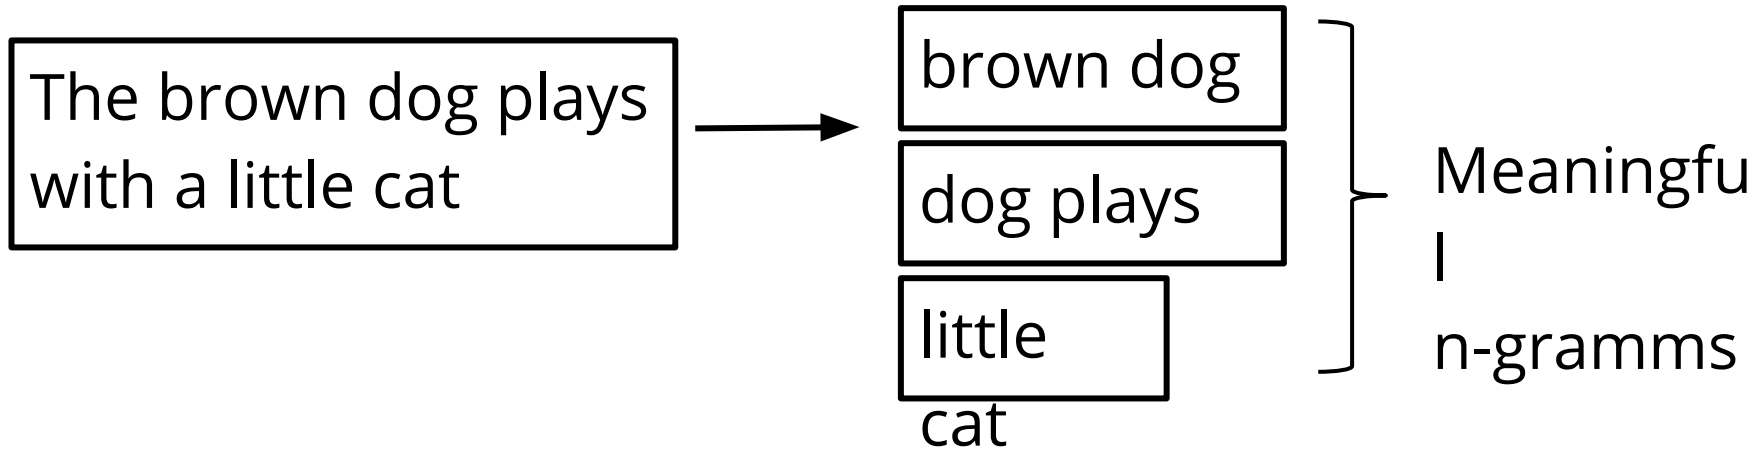
a little

little

cat



Do we need all this bigramms?



Meaningful n-grams are often called **collocations**

How to detect meaningful n-grams?



Delete:

- ▷ High-frequency n-gramms
 - Articles, prepositions
 - Auxiliary verbs (to be, to have, etc.)
 - General vocabulary
- ▷ Low-frequency n-gramms
 - Typos
 - Combinations that occur 1-2 times in a text



Term Frequency (tf): gives us the frequency of the word in each document in the corpus.

$$\text{tf}(t, d) = f_{t, d}$$



Inverse Document Frequency (idf): used to calculate the weight of rare words across all documents in the corpus. The words that occur rarely in the corpus have a high IDF score.

$$\text{idf}(t, D) = \log \frac{N}{|\{d \in D : t \in d\}|}$$

N : total number of documents in the corpus $N = |D|$

$|\{d \in D : t \in d\}|$: number of documents where the term t appears



Sentence A: The car is driven on the road.



Sentence B: The truck is driven on the highway.

(each sentence is a separate document)



Word	TF		IDF	TF * IDF	
	A	B		A	B
The	1/7	1/7			
Car	1/7	0			
Truck	0	1/7			
Is	1/7	1/7			
Driven	1/7	1/7			
On	1/7	1/7			
The	1/7	1/7			
Road	1/7	0			
Highway	0	1/7			



Word	TF		IDF	TF * IDF	
	A	B		A	B
The	1/7	1/7	$\log(2/2)=0$		
Car	1/7	0	$\log(2/1)=0.3$		
Truck	0	1/7	$\log(2/1)=0.3$		
Is	1/7	1/7	$\log(2/2)=0$		
Driven	1/7	1/7	$\log(2/2)=0$		
On	1/7	1/7	$\log(2/2)=0$		
The	1/7	1/7	$\log(2/2)=0$		
Road	1/7	0	$\log(2/1)=0.3$		
Highway	0	1/7	$\log(2/1)=0.3$		



Word	TF		IDF	TF * IDF	
	A	B		A	B
The	1/7	1/7	$\log(2/2)=0$	0	0
Car	1/7	0	$\log(2/1)=0.3$	0.043	0
Truck	0	1/7	$\log(2/1)=0.3$	0	0.043
Is	1/7	1/7	$\log(2/2)=0$	0	0
Driven	1/7	1/7	$\log(2/2)=0$	0	0
On	1/7	1/7	$\log(2/2)=0$	0	0
The	1/7	1/7	$\log(2/2)=0$	0	0
Road	1/7	0	$\log(2/1)=0.3$	0.043	0
Highway	0	1/7	$\log(2/1)=0.3$	0	0.043



```
from sklearn.feature_extraction.text  
  
import TfidfVectorizer
```





Word Embeddings



One-hot vectors:

One-hot vectors

Rome = [1, 0, 0, 0, 0, 0, ..., 0]

Paris = [0, 1, 0, 0, 0, 0, ..., 0]

Italy = [0, 0, 1, 0, 0, 0, ..., 0]

France = [0, 0, 0, 1, 0, 0, ..., 0]

word V

Problems:



Huge vectors



VERY sparse



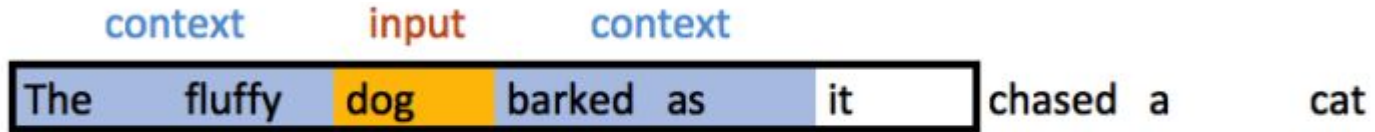
No semantics or word similarity information included



Does vector similarity imply semantic similarity?

"You shall know a word by the company it keeps"

Firth, 1957






Word representations via matrix factorization



Input: PMI, word cooccurrences, etc.

Method: dimensionality reduction (SVD)

Output: word similarities





Delete:

- ▷ High-frequency n-gramms
 - Articles, prepositions
 - Auxiliary verbs (to be, to have, etc.)
 - General vocabulary
- ▷ Low-frequency n-gramms
 - Typos
 - Combinations that occur 1-2 times in a text



Collocations: context is all you need

- ▶ Cooccurrence counters in a window of fixed size
 - ▶ n_{uv} states for the number of times we've seen word u and word v together in the window
- Better solution: Pointwise Mutual Information (PMI)

$$PMI = \log \frac{p(u, v)}{p(u)p(v)} = \log \frac{n_{uv}n}{n_u n_v}$$

- Much better solution: **Positive PMI (pPMI)**

$$pPMI = \max(0, PMI)$$

Frequency With Filter		PMI	T-test With Filter	Chi-Sq Test
(front, desk)	(universal, studios)		(front, desk)	(wi, fi)
(great, location)	(howard, johnson)		(great, location)	(cracker, barrel)
(friendly, staff)	(cracker, barrel)		(friendly, staff)	(howard, johnson)
(hot, tub)	(santa, barbara)		(hot, tub)	(la, quinta)
(clean, room)	(sub, par)	(continental, breakfast)		(front, desk)
(hotel, staff)	(santana, row)	(free, breakfast)		(universal, studios)
(continental, breakfast)	(e, g)	(great, place)		(santa, barbara)
(nice, hotel)	(elk, springs)	(parking, lot)		(santana, row)
(free, breakfast)	(times, square)	(customer, service)		(, more)
(great, place)	(ear, plug)	(desk, staff)		(flat, screen)
(desk, staff)	(la, quinta)	(walk, distance)		(french, quarter)
(parking, lot)	(fire, pit)	(comfortable, bed)		(elk, springs)
(customer, service)	(san, clemente)	(nice, hotel)		(walking, distance)




Why not to learn word vectors?



Embeddings: intuition


What is $\text{king} - \text{man} + \text{woman}$?






Embeddings: intuition

So $\text{king} - \text{man} + \text{woman} = \text{queen!}$







Word2vec (Mikolov et al. 2013) - a framework for learning word embeddings



Male-Female



Verb tense



Country-Capital

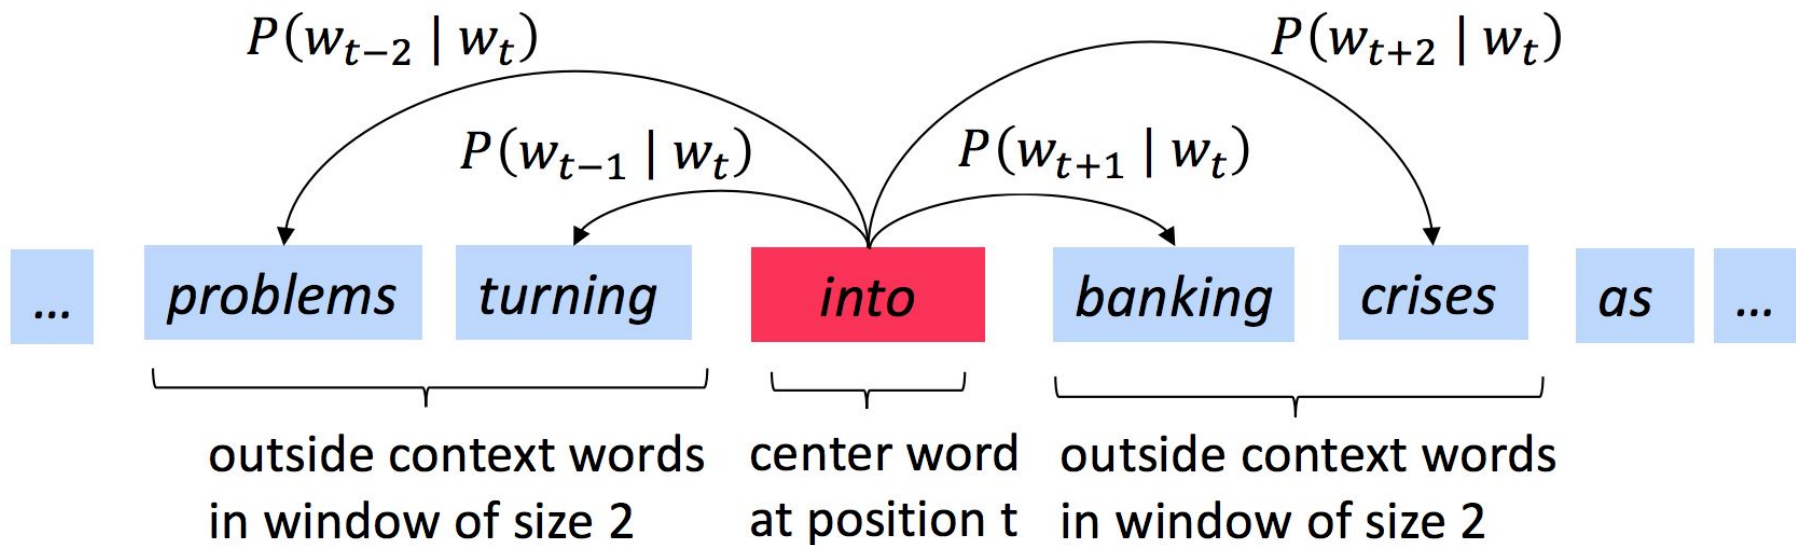


Embeddings: word2vec

Source Text

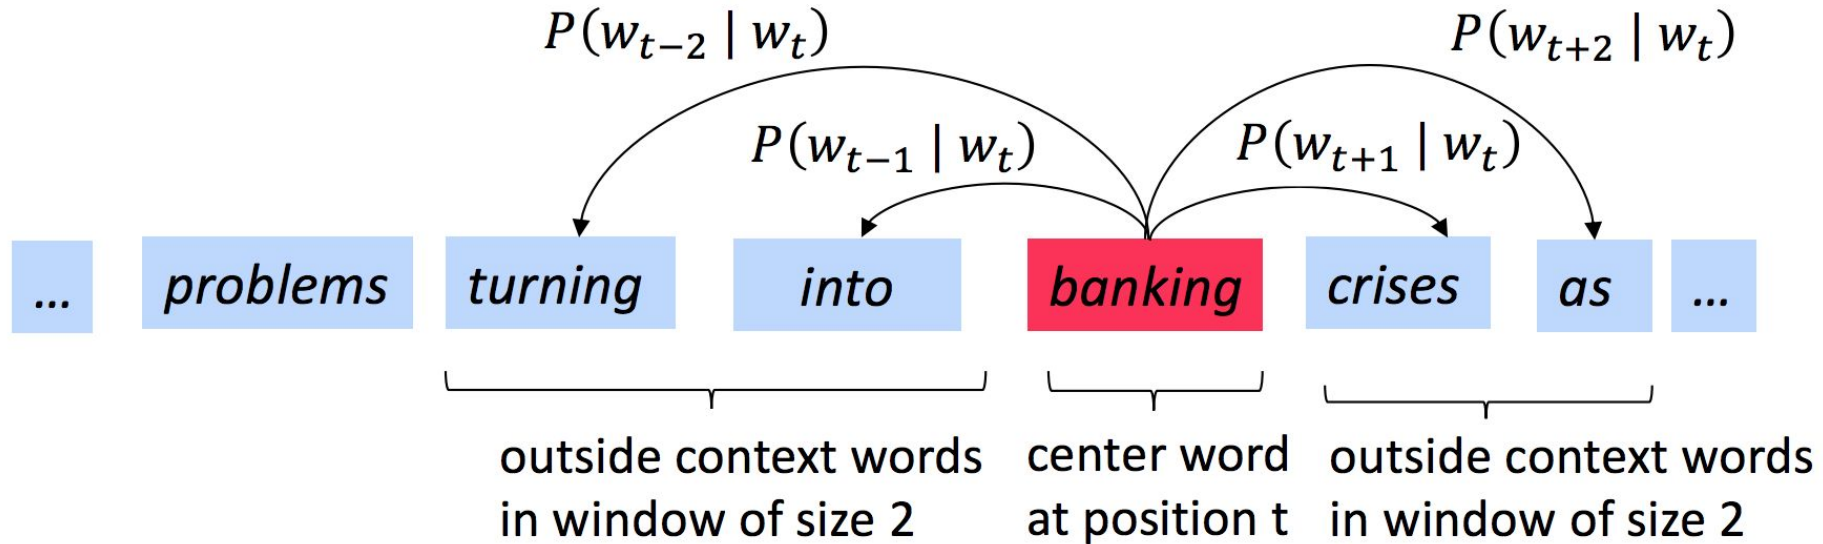
Training Samples

The quick brown fox jumps over the lazy dog. →	(the, quick) (the, brown)
The quick brown fox jumps over the lazy dog. →	(quick, the) (quick, brown) (quick, fox)
The quick brown fox jumps over the lazy dog. →	(brown, the) (brown, quick) (brown, fox) (brown, jumps)
The quick brown fox jumps over the lazy dog. →	(fox, quick) (fox, brown) (fox, jumps) (fox, over)



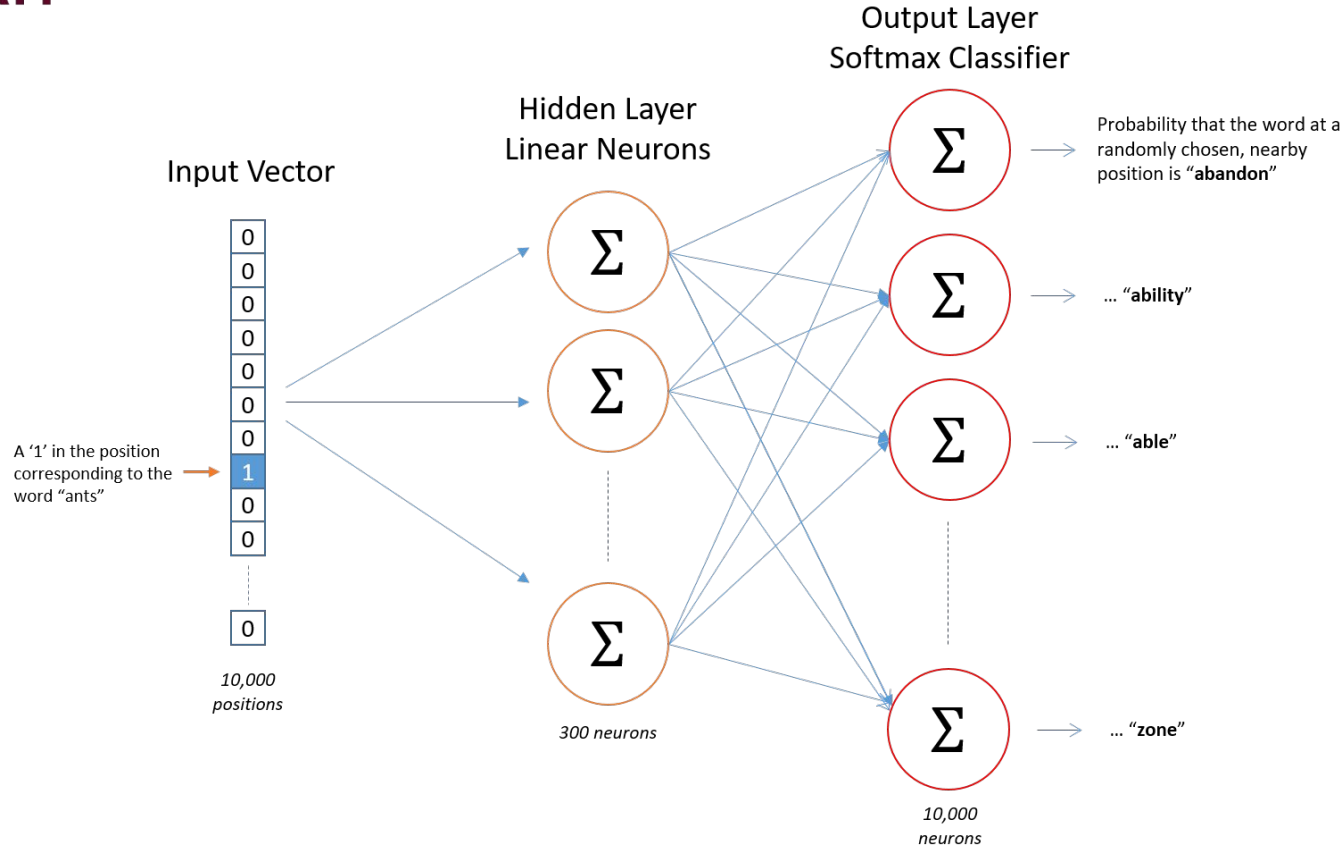


Embeddings: word2vec






Embeddings: word2vec





Embeddings: word2vec





- ▶ Word vectors with 300 components
- ▶ Vocabulary of 10,000 words.
- ▶ Weight matrix with $300 \times 10,000 = 3$ million weights each!

Training is too long and computationally expensive
How to fix this?


Basic approaches:

1. Treating common word pairs or phrases as single “words” in their model.
2. Subsampling frequent words to decrease the number of training examples.
3. Modifying the optimization objective with a technique they called “Negative Sampling”, which causes each training sample to update only a small percentage of the model’s weights.

Subsampling frequent words.

w_i is the word, $z(w_i)$ is the fraction of this word in the text

Graph for $(\sqrt{x/0.001}+1)*0.001/x$



$P(w_i)$ is the probability of *keeping* the word:

$$P(w_i) = \left(\sqrt{\frac{z(w_i)}{0.001}} + 1 \right) \cdot \frac{0.001}{z(w_i)}$$



Embeddings: negative sampling

Negative Sampling idea: only few words error is computed. All other words have zero error, so no updates by the backprop mechanism.

More frequent words are selected to be negative samples more often. The probability for selecting a word is just its weight divided by the sum of weights for all words.

$$P(w_i) = \frac{f(w_i)^{3/4}}{\sum_{j=0}^n (f(w_j)^{3/4})}$$



Continuous BOW (CBOW)

$$p(w_i | w_{i-h}, \dots, w_{i+h})$$

Predict center word from
(bag of) context words


- ▶ Predicting one word each time
- ▶ Relatively fast


Skip-gram


$$p(w_{i-h}, \dots, w_{i+h} | w_i)$$

Predict context ("outside")
words (position independent)
given center word

- ▶ Predicting context by one word
- ▶ Much slower
- ▶ Better with infrequent words










GloVe Visualizations: Company - CEO







Word2vec: word analogies

King - man + woman = queen
↓ ↓ ↓ ↓
 x y y' $target$

$\cos(x - y + y', target) \rightarrow \max_{target}$





Word vectors are simply vectors of numbers that represent the meaning of a word

Approaches:

- ▶ One-hot encoding
- ▶ Bag-of-words models
- ▶ Counts of word / context co-occurrences
- ▶ TF-IDF
- ▶ Predictions of context given word (skip-gram neural network models, e.g. word2vec)



**BIGDATA
TEAM**

Backup